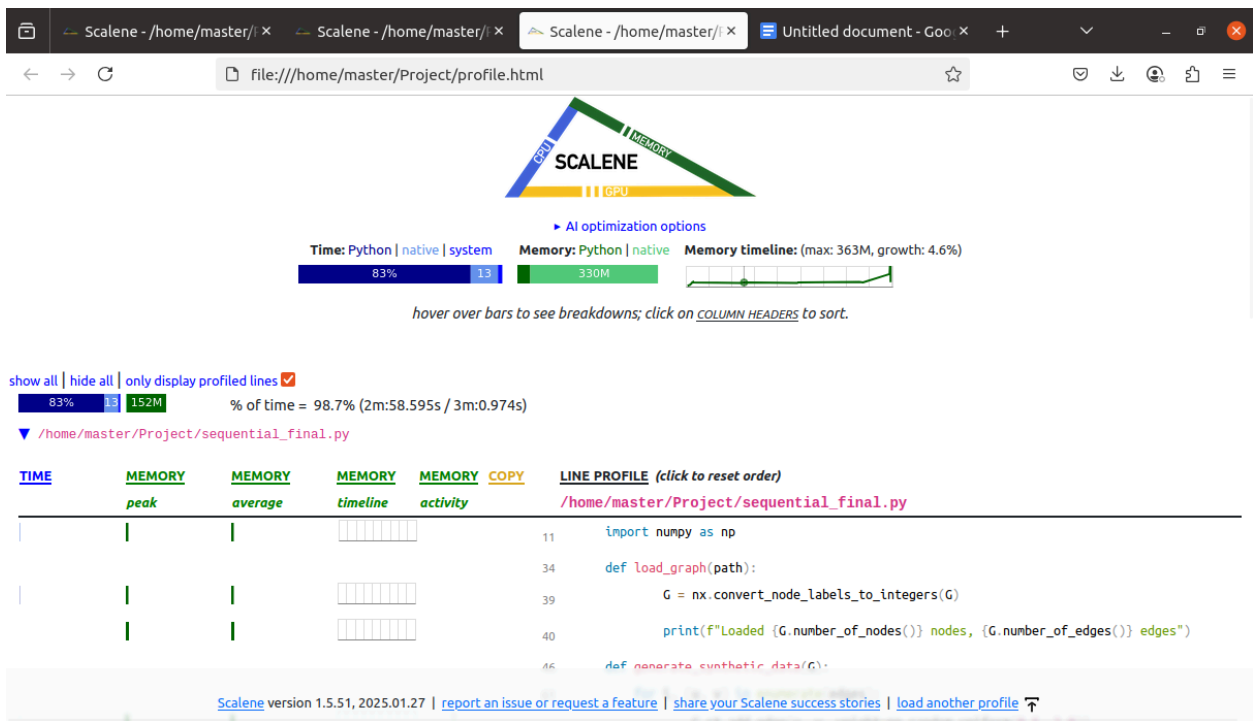
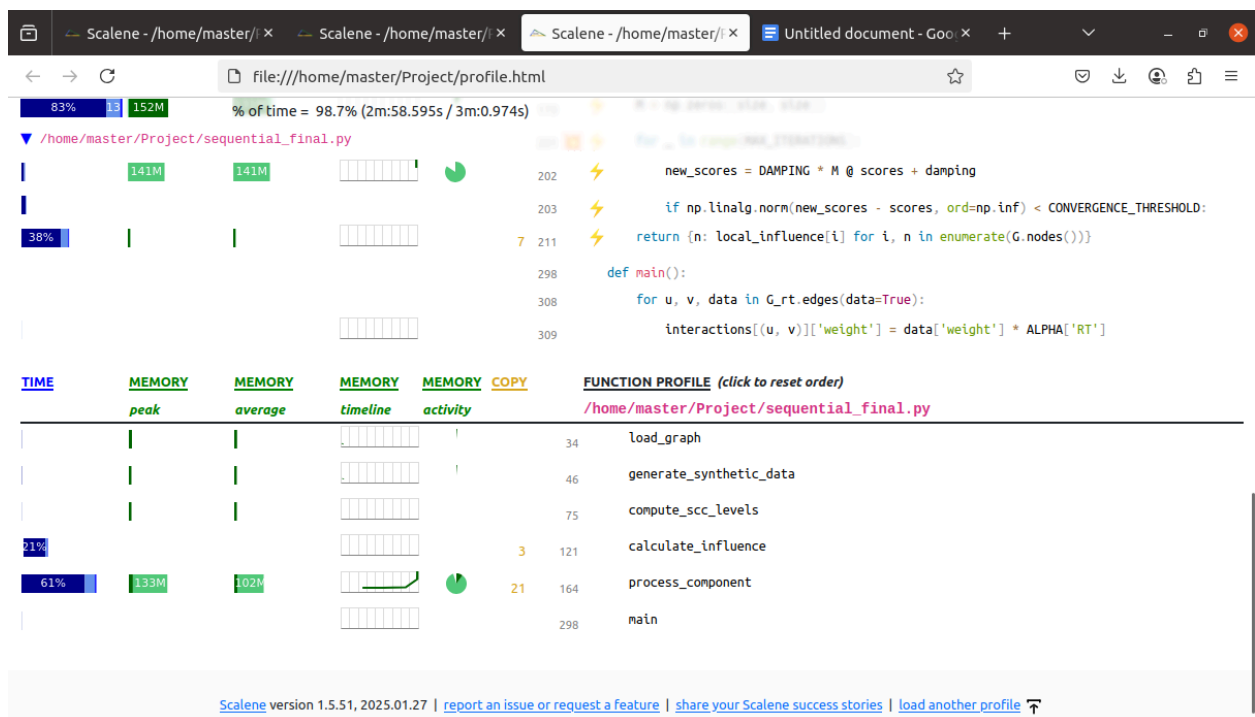
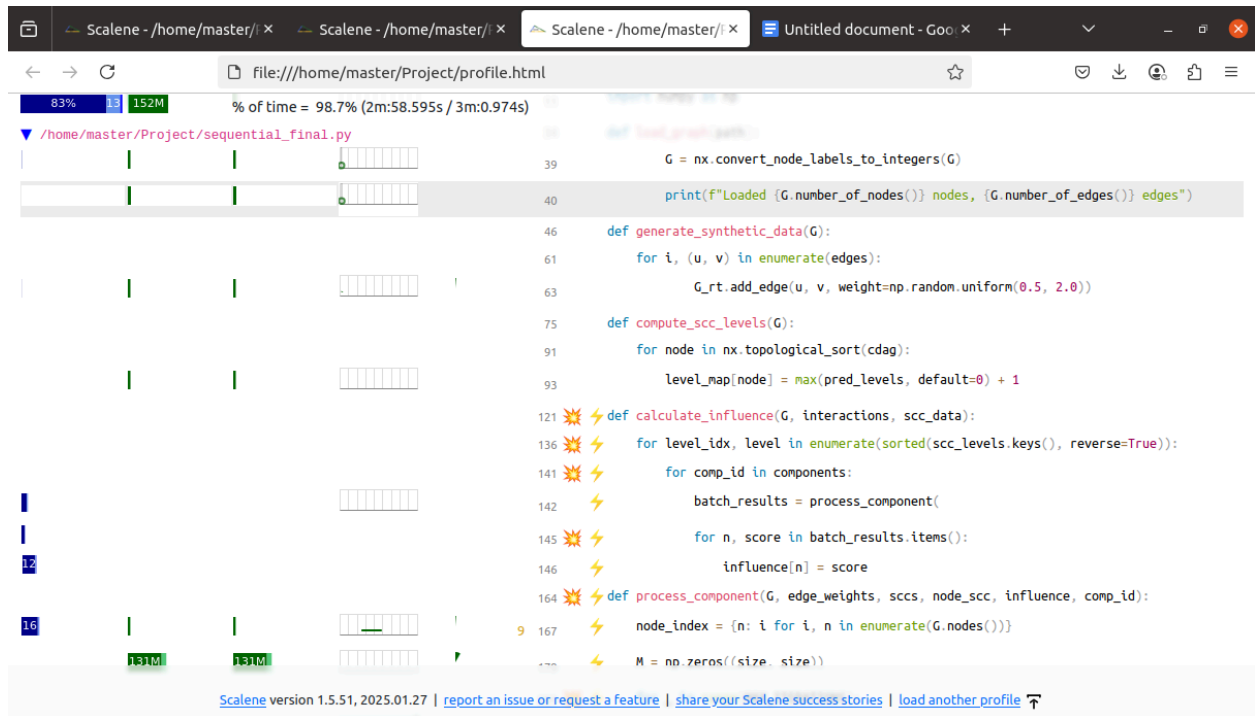


PDC PROJECT REPORT

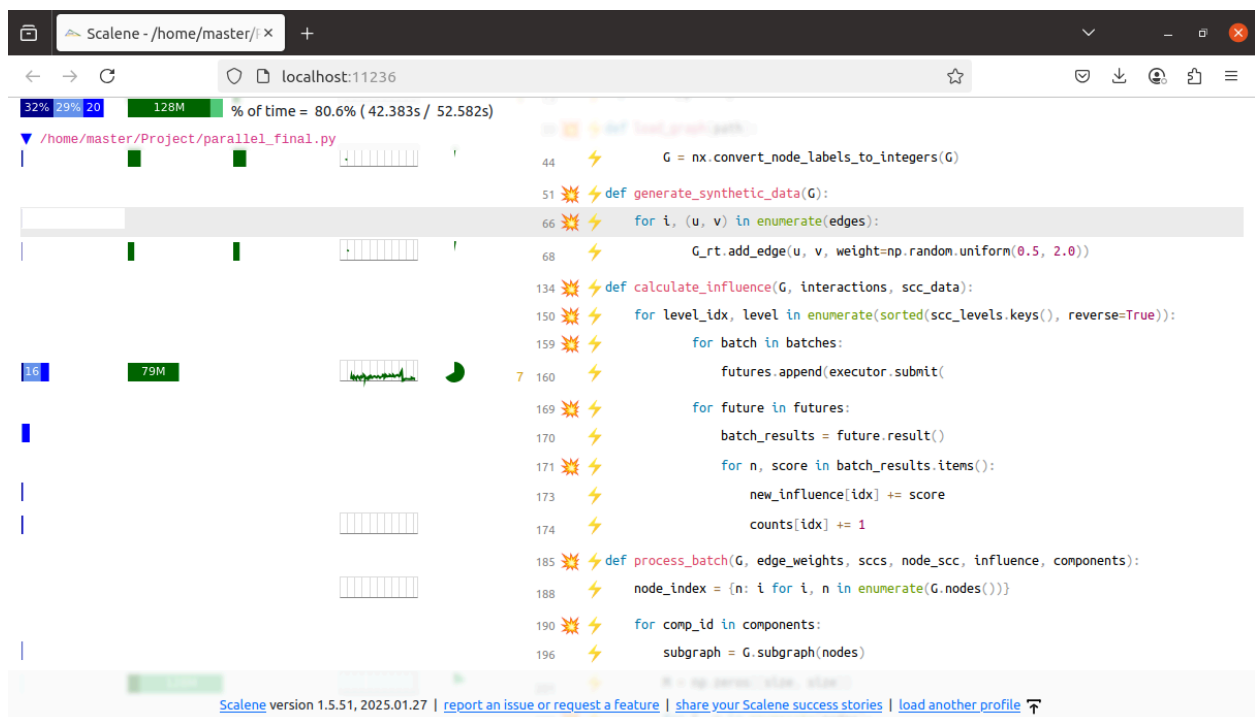
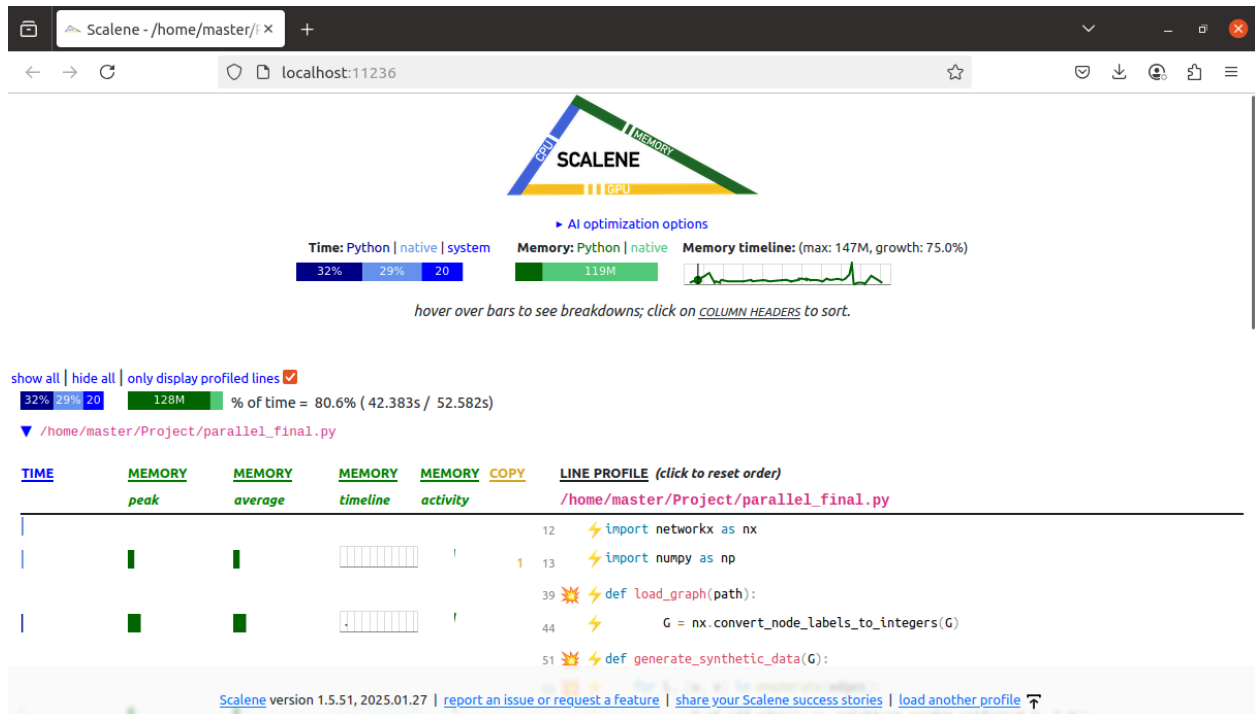
Scalene Analysis (Best for deep analysis)
Used for all 3 (Time + memory + I/O)

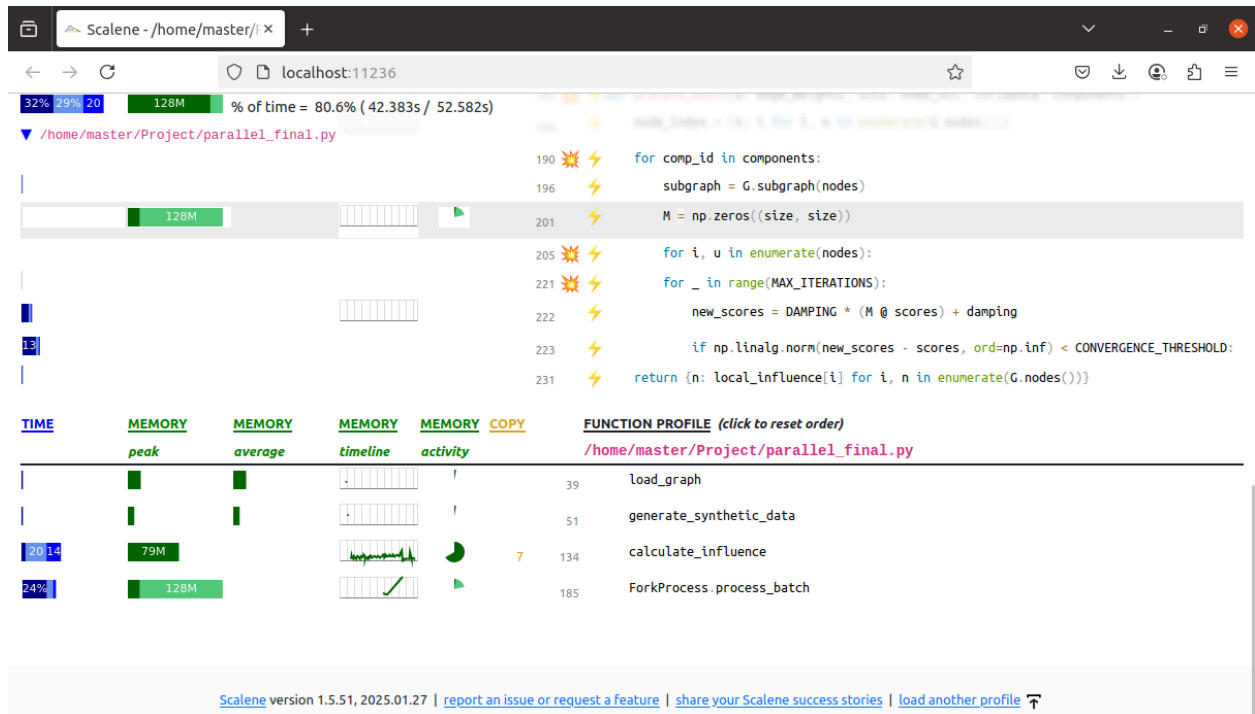
Sequential:





Parallel:





Time comparison:

Sequential : Approx 114 sec

Parallel: Approx 22 sec

MPI : Approx 12 sec

MPI+(OpenMP) : Approx 10 sec

Speedup:

Speedup = Sequential Time / Parallel Time

Speedup Parallel = $114 / 22 \approx 5.18$

Speedup MPI = $114 / 12 \approx 9.5$

Speedup for MPI + OpenMP (Hybrid) $\approx 114 / 10 \approx 11.4$

Code Explanaton:

1) file(Sequential.py) :

This Python script implements a **sequential influence maximization algorithm** (called PSALIM) on a **peer-to-peer (P2P) network graph** from the `p2p-Gnutella04` dataset. It removes parallelism to keep things simple and focuses on understanding how influence spreads in a social-like network using synthetic interaction data.

◆ Step-by-Step Breakdown:

1. Graph Loading:

The script loads a directed graph from a text file (`p2p-Gnutella04.txt`). Each line represents an edge, indicating a connection from one node to another.

2. Synthetic Interaction Generation:

To simulate real-world social interactions, the script generates **three types of interactions** randomly over the network edges:

- **Retweet (RT)** – 30% chance
 - **Reply (RE)** – 20% chance
 - **Mention (MT)** – 10% chance
- Each interaction is given a random weight to reflect how strong or frequent it is.

3. Community Detection (SCC Levels):

The graph is analyzed to find **Strongly Connected Components (SCCs)**. These are clusters of nodes where each node can reach every other node in the same component.

These components are organized into **levels** using a **condensed DAG (Directed Acyclic Graph)** formed by collapsing SCCs into single nodes, and then topologically sorting them. This helps in processing influence from higher (source) communities down to lower ones.

4. Influence Calculation:

The core idea is that nodes in a network influence each other based on the structure and the interactions. For each SCC:

- A **transition matrix** is built to model how influence flows between nodes based on edge weights.
- Then, **power iteration** is used to compute influence scores for the nodes in each SCC.
- The influence is **updated level by level**, ensuring higher-level nodes propagate their influence downward.

5. Most Influential Nodes per Community:

For the **top 10 largest communities**, the script identifies the **top 3 most influential nodes** based on the computed influence scores.

6. Seed Selection:

The algorithm selects a set of **top-k nodes (default 10)** to act as “**seeds**” for influence spread. It uses a **community-aware greedy strategy**, considering:

- High influence scores globally
- Top nodes within the largest communities
This ensures that the selected nodes are both influential and well-distributed across the network structure.

2) file(parallel.py) :

This file aims to:

1. **Analyze the influence of nodes** in a P2P network graph.
2. **Detect communities (SCCs)** in the graph.
3. **Calculate influence scores** within and across these communities.
4. **Select seed nodes** with high influence for diffusion or propagation tasks (like information spread).

Overall Workflow

- 3) **Load Graph:** Reads a `.txt` file representing the directed graph (`p2p-Gnutella04.txt`) using NetworkX.
- 4) **Generate Interaction Types:**
 - a) Randomly simulates 3 types of interactions on the edges:

- i) Retweets (RT)
- ii) Replies (RE)
- iii) Mentions (MT)

b) Assigns weights to these interactions for influence modeling.

5) **Community Detection via SCCs:**

- a) Computes **Strongly Connected Components (SCCs)** of the graph.
- b) Converts them into a **condensed DAG** to model levels (hierarchies of communities).
- c) Each SCC is assigned a "level" based on topological sorting.

6) **Influence Computation:**

- a) Influence scores are initialized with slight randomness and normalized.
- b) The algorithm goes through SCC levels in reverse order.
- c) For each SCC, it calculates influence via **power iteration** using a transition matrix (similar to PageRank).
- d) This is done **in parallel using multiple CPU cores** for efficiency.

7) **Show Influential Nodes:**

- a) From the largest communities, it selects and prints the **top 3 most influential nodes** in each using their computed scores.
- 8) **Seed Selection:**
 - a) Picks **k** (e.g., 10) nodes with the **highest influence scores**, distributed across different communities, ensuring diversity.

3) file (mpi.py):

What This Code Accomplishes:

- **Distributed Graph Processing:** Each MPI rank loads and partitions the graph, and works on a subgraph.
- **Threaded Processing:** Within each MPI process, edge operations and influence calculations are parallelized with threads.
- **Synthetic Modeling:** Three synthetic edge types simulate different interaction modes (RT/RE/MT).
- **PageRank-Like Influence Scoring:** Uses a weighted version of PageRank within each SCC level.
- **Seed Selection:** Top-k influential nodes per partition are selected and gathered at root.

Suggestions for Improvement or Extension:

- 9) **File Loading Optimization:**

- a) All nodes currently load the full graph (`G_full = load_graph(...)`). This could be memory-inefficient for large graphs.
- b) **Alternative:** Only rank 0 loads and partitions, then scatters subgraphs to other ranks.

10) Edge Storage:

- a) Edge weights are recomputed every time from scratch.
- b) **Optimization:** Persist or cache synthetic edge graphs and precomputed weights to avoid recomputation on each run.

11) Influence Propagation Between Partitions:

- a) Influence calculation is done per partition only.
- b) **Limitation:** Inter-partition influence is not accounted for (important if cross-partition edges exist).
- c) **Fix:** Implement **ghost nodes** or **boundary synchronization** between partitions.

12) Logging Improvements:

- a) Current `print()` usage is clean but mixed.
- b) You could consider using the `logging` module with rank-specific prefixes for consistency and better control.

13) Scalability Analysis:

- a) Include code for measuring:
 - i) Speedup per MPI process count
 - ii) Breakdown of computation vs communication time

14) Testing on p2p-Gnutella04:

- a) Ensure file format matches expected edge list (directed). You may want to validate the dataset shape before full runs.

4) file (mpi_open-mp.py):

This program performs influence maximization on a graph (like a social network) using parallel processing. It identifies the most "influential" nodes in the network, i.e., the best starting points to spread information.

Key Technologies Used

- MPI (**mpi4py**): Distributes work across multiple processes (machines or cores).
- OpenMP-style threading (**ThreadPoolExecutor**): Uses threads for faster processing within each MPI process.
- METIS: Optional, for smart graph partitioning.
- NetworkX: For graph operations (loading, SCCs, etc.).

Step-by-Step Breakdown

1. Initialization

- Uses MPI to determine which process you are (rank).
- Uses available CPU cores to decide how many threads to run.

2. Graph Loading

- Loads a real-world graph (`p2p-Gnutella04.txt`) into memory using NetworkX.
- All MPI processes load the full graph.

3. Graph Partitioning

- If METIS is available, partitions the graph evenly across MPI processes for load balancing.
- Otherwise, it splits nodes manually based on index.

4. Synthetic Data Generation

- Creates three interaction graphs (`RT`, `RE`, `MT`) with randomly weighted edges.
- This simulates different types of relationships in the network.

5. Community Detection

- Computes Strongly Connected Components (SCCs).
- Organizes these SCCs into levels, where level 1 has no incoming edges from other SCCs.

6. Edge Weight Computation

- Precomputes edge weights by combining the weights from RT, RE, and MT graphs using predefined importance factors (ALPHA).

7. Influence Calculation

- Performs influence propagation per SCC level using a matrix-based iterative method (like PageRank).
- Processes SCCs from highest to lowest level to respect the direction of influence.

8. Seed Node Selection

- Picks top $k=10$ nodes with the highest influence scores within each MPI process.

9. Result Gathering

- Master node (rank 0) gathers all top seeds from workers and prints them.

10. Execution Time

- Prints total execution time at the end.
-